

# Biostatistical Methods I Final Project

Benjamin Goebel (bpg2118), Jesse Ames (jra2186), Louis Sharp (lzs2109), Brian Lee (jl6046)

12/17/2021

Let's define helpful functions.

```
# Functions

# Purpose: Calculates the Pearson's correlation coefficient between every
# variable in the data set and a specified variable.
# Arguments: v_name: a variable of type character that is the variable name
# Returns: A knitted table of correlations.
get_cor_by_var <- function(v_name) {
  cdi_slim %>%
  map(~cor(as.numeric(.x), pull(cdi, v_name), method = "pearson")) %>%
  as_tibble() %>%
  pivot_longer(CRM_1000:log_pop_density,
               names_to = "variables",
               values_to = "r") %>%

  mutate(
    sign = ifelse(r < 0, "-", "+"),
    r = abs(r)
  ) %>%
  arrange(desc(r)) %>%
  knitr::kable()
}

# Purpose: Fits the model and gets the model adjusted r-squared.
# Arguments: mod: a variable of type character that is the formula to fit a
# linear model.
# Returns: A numeric, the model adjusted r-squared.
get_mod_adj_r_squared <- function(mod, data = cdi) {
  lm(mod, data = data) %>%
  broom::glance() %>%
  pull(adj.r.squared)
}

# Purpose: Performs cross validation on a model specified by its formula.
# Arguments: mod: a variable of type character that is the formula to fit a
# linear model.
# Returns: A column vector of the model root mean squared errors generated by the validation procedure.
get_cv_rmse <- function(mod, data = cdi) {
  set.seed(1)
  crossv_mc(data, 1000) %>%
  mutate(
    train = map(train, as_tibble),

```

```

    test = map(test, as_tibble)
  ) %>%
  mutate(
    fitted_mod = map(train, ~lm(mod, data = .x))
  ) %>%
  mutate(
    rmse_mod = map2_dbl(fitted_mod, test, ~rmse(model = .x, data = .y))
  ) %>%
  pull(rmse_mod)
}

# Purpose: Fits model.
#           Plots model residual as a function of model prediction for the given
#           model formula.
# Arguments: mod: a variable of type character that is the formula to fit a
#           linear model.
# Returns: The ggplot.
plot_model_residuals <- function(mod, data = cdi) {
  fitted_mod <- lm(mod, data = data)
  cdi %>%
  add_predictions(fitted_mod) %>%
  add_residuals(fitted_mod) %>%
  ggplot(aes(x = pred, y = resid)) +
  geom_point() +
  theme_bw() +
  labs(
    title = ""
  ) +
  theme(plot.title = element_text(hjust = 0.5))
}

# Purpose: Fits the specified model and creates a QQ Plot.
# Arguments: mod: a variable of type character that is the formula to fit a
#           linear model.
# Returns: The plot.
plot_mod_qq <- function(mod, data = cdi) {
  mod %>%
  lm(data = data) %>%
  plot(which = 2)
}

# Purpose: Fits the specified model and creates a leverage plot.
# Arguments: mod: a variable of type character that is the formula to fit a
#           linear model.
# Returns: The plot.
plot_mod_leverage <- function(mod, data = cdi) {
  mod %>%
  lm(data = data) %>%
  plot(which = 5)
}

```

Let's begin by reading in the data and adding a column for the crime rate per 1,000 people in the county

population. We will name this column CRM\_1000. We will then recode the region variable as a factor.

```
cdi <- read_csv(here::here("data", "cdi.csv")) %>%
  mutate(CRM_1000 = (crimes/pop) * 1000,
         state = as.factor(state),
         region = as.factor(region),
         region = fct_recode(region, "Northeast" = "1", "North Central" = "2",
                             "South" = "3", "West" = "4"),
         pop_density = pop/area,
         log_pop18 = log(pop18),
         log_pop65 = log(pop65),
         log_hsgrad = log(hsgrad),
         log_bagrads = log(bagrads),
         log_poverty = log(poverty),
         log_unemp = log(unemp),
         log_totalinc = log(totalinc),
         log_pcincome = log(pcincome),
         log_pop_density = log(pop_density)
        ) %>%
  dplyr::select(-id, -state, -cty, -docs, -beds, -crimes, -pop, -area) %>%
  dplyr::select(CRM_1000, region, everything())
```

Let's compare each variable's distribution in the cdi dataset as is and log-transformed to see which is more normal.

```
cdi_long =
  cdi %>%
    pivot_longer(
      pop18:log_pop_density,
      names_to = "var",
      values_to = "val"
    ) %>%
  mutate(
    trans = case_when(grepl('log_', var) ~ 'log',
                     TRUE ~ 'ori'
    ),
    trans = factor(trans, levels = c("ori", "log"), labels = c("Untransformed", "log Transformed")),
    var = str_replace(var, "log_", ""),
    var = factor(var, levels = c("pop18", "pop65", "hsgrad", "bagrad", "poverty", "unemp", "pcincome",
                                "CRM_1000", "region"), labels = c("pop18", "pop65", "hsgrad", "bagrad", "poverty", "unemp", "pcincome",
                                                                    "CRM_1000", "region"))
  )

dens_plot_gen = function(v, b){

  if(!b){
    cdi_long %>%
      filter(var == v) %>%
      ggplot(aes(x = val)) +
      geom_boxplot() +
      theme_minimal() +
      theme(
        axis.title.y=element_blank(),
        axis.title.x=element_blank(),
        axis.text.y=element_blank(),
        axis.text.x=element_blank()
      )
  }
}
```

```

) +
facet_grid(var ~ trans, switch = "y", scales = "free") +
theme(
  strip.text.y.left = element_text(angle = 0),
  strip.text.x = element_blank()
)
}else{
  cdi_long %>%
  filter(var == v) %>%
  ggplot(aes(x = val)) +
  geom_boxplot() +
  theme_minimal() +
  theme(
    axis.title.y=element_blank(),
    axis.title.x=element_blank(),
    axis.text.y=element_blank(),
    axis.text.x=element_blank()
  ) +
  facet_grid(var ~ trans, switch = "y", scales = "free") +
  theme(
    strip.text.y.left = element_text(angle = 0)
  )
}
}

var_list = c("pop18", "pop65", "hsgrad", "bagrads", "poverty", "unemp", "pcincome", "totalinc", "pop_dens")

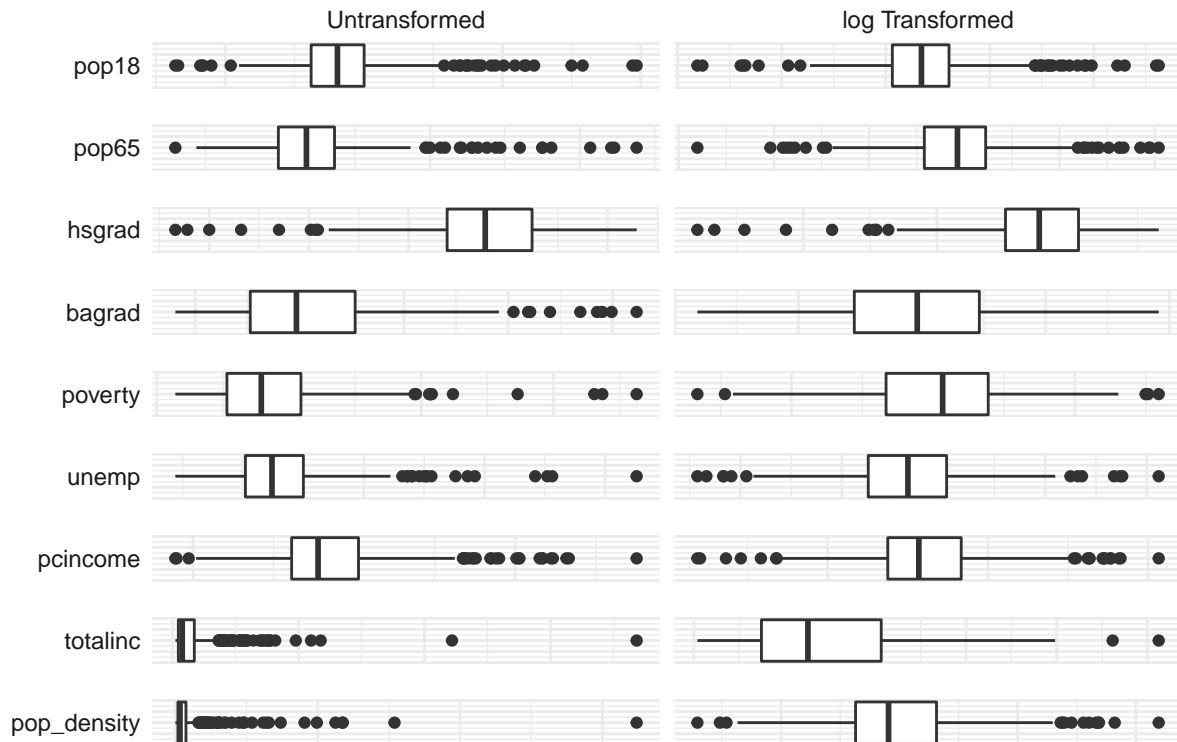
plot_1 = dens_plot_gen(var_list[1], T)
plot_2 = dens_plot_gen(var_list[2], F)
plot_3 = dens_plot_gen(var_list[3], F)
plot_4 = dens_plot_gen(var_list[4], F)
plot_5 = dens_plot_gen(var_list[5], F)
plot_6 = dens_plot_gen(var_list[6], F)
plot_7 = dens_plot_gen(var_list[7], F)
plot_8 = dens_plot_gen(var_list[8], F)
plot_9 = dens_plot_gen(var_list[9], F)

patch_1 = plot_1 / plot_2 / plot_3 / plot_4 / plot_5 / plot_6 / plot_7 / plot_8 / plot_9

patch_1 +
  plot_annotation(title = 'Untransformed and log Transformed Variable Distributions in the CDI Dataset',
    theme = theme(plot.title = element_text(size = 14)))

```

## Untransformed and log Transformed Variable Distributions in the CDI Data



Based on this boxplot, we decided to log transform all continuous variables except high school graduation. We can remove all unused variables now.

```
cdi_slim <-
  cdi %>%
  select(-c(pop18, pop65, bagrad, poverty, unemp, pcincome, totalinc,
             pop_density, log_hsgrad))
```

Now, let's calculate the Pearson's correlation coefficient between every variable in the data set and CRM\_1000.

```
get_cor_by_var("CRM_1000")
```

| variables       | r         | sign |
|-----------------|-----------|------|
| CRM_1000        | 1.0000000 | +    |
| log_poverty     | 0.4823623 | +    |
| region          | 0.3427584 | +    |
| log_pop_density | 0.3367361 | +    |
| log_totalinc    | 0.3273042 | +    |
| hsgrad          | 0.2264129 | -    |
| log_pop18       | 0.2039079 | +    |
| log_pcincome    | 0.0695287 | -    |
| log_bagrad      | 0.0632119 | +    |
| log_pop65       | 0.0543376 | -    |
| log_unemp       | 0.0362602 | +    |

Next, we can perform stepwise selection based on AIC.

```
# keep only the predictors needed to build the model using pcincome as income measure
pcincome_model =
  cdi %>%
  dplyr::select(-c(pop18, pop65, bagrad, poverty, unemp, pcincome, totalinc, log_totalinc, pop_density,

# get first full mlr
fit_1 = lm(CRM_1000 ~ ., data = pcincome_model)
## summary(fit_1) # aRs = 0.5309
## boxcox(fit_1) # we ignore the square root tranformation for a more sensible interpretation

# run stepwise and get a list of highly effective predictors
step(fit_1, direction = 'both')
```

```
## Start:  AIC=2588.89
## CRM_1000 ~ region + hsgrad + log_pop18 + log_pop65 + log_bagrad +
##      log_poverty + log_unemp + log_pcincome + log_pop_density
##
##              Df Sum of Sq    RSS    AIC
## - hsgrad      1      7.3 149659 2586.9
## - log_pop65    1     21.0 149672 2586.9
## - log_bagrad   1     41.6 149693 2587.0
## - log_unemp    1    150.4 149802 2587.3
## <none>                149651 2588.9
## - log_pop18    1    860.1 150512 2589.4
## - log_pcincome  1   4505.8 154157 2599.9
## - log_pop_density 1  18265.0 167916 2637.6
## - region       3  24405.2 174057 2649.4
## - log_poverty  1  30268.1 179920 2667.9
##
## Step:  AIC=2586.91
## CRM_1000 ~ region + log_pop18 + log_pop65 + log_bagrad + log_poverty +
##      log_unemp + log_pcincome + log_pop_density
##
##              Df Sum of Sq    RSS    AIC
## - log_pop65    1      21 149680 2585.0
## - log_bagrad   1      35 149694 2585.0
## - log_unemp    1     164 149823 2585.4
## <none>                149659 2586.9
## - log_pop18    1     875 150534 2587.5
## + hsgrad       1       7 149651 2588.9
## - log_pcincome  1    4702 154361 2598.5
## - log_pop_density 1  18484 168143 2636.2
## - region       3   24678 174337 2648.1
## - log_poverty  1   38495 188154 2685.6
##
## Step:  AIC=2584.97
## CRM_1000 ~ region + log_pop18 + log_bagrad + log_poverty + log_unemp +
##      log_pcincome + log_pop_density
##
##              Df Sum of Sq    RSS    AIC
## - log_bagrad   1      42 149721 2583.1
## - log_unemp    1     174 149853 2583.5
```

```

## <none> 149680 2585.0
## + log_pop65 1 21 149659 2586.9
## + hsgrad 1 7 149672 2586.9
## - log_pop18 1 1472 151152 2587.3
## - log_pcincome 1 4829 154509 2596.9
## - log_pop_density 1 18551 168230 2634.4
## - region 3 27888 177568 2654.2
## - log_poverty 1 44711 194390 2698.0
##
## Step: AIC=2583.09
## CRM_1000 ~ region + log_pop18 + log_poverty + log_unemp + log_pcincome +
## log_pop_density
##
## Df Sum of Sq RSS AIC
## - log_unemp 1 132 149854 2581.5
## <none> 149721 2583.1
## + log_bagrad 1 42 149680 2585.0
## + log_pop65 1 27 149694 2585.0
## + hsgrad 1 1 149720 2585.1
## - log_pop18 1 2731 152452 2589.1
## - log_pcincome 1 8413 158134 2605.2
## - log_pop_density 1 18534 168255 2632.4
## - region 3 29058 178780 2655.1
## - log_poverty 1 44988 194710 2696.7
##
## Step: AIC=2581.48
## CRM_1000 ~ region + log_pop18 + log_poverty + log_pcincome +
## log_pop_density
##
## Df Sum of Sq RSS AIC
## <none> 149854 2581.5
## + log_unemp 1 132 149721 2583.1
## + log_pop65 1 31 149823 2583.4
## + hsgrad 1 10 149844 2583.4
## + log_bagrad 1 1 149853 2583.5
## - log_pop18 1 2657 152511 2587.2
## - log_pcincome 1 8281 158135 2603.2
## - log_pop_density 1 18596 168450 2630.9
## - region 3 31014 180868 2658.2
## - log_poverty 1 51449 201303 2709.3
##
##
## Call:
## lm(formula = CRM_1000 ~ region + log_pop18 + log_poverty + log_pcincome +
## log_pop_density, data = pcincome_model)
##
## Coefficients:
## (Intercept) regionNorth Central regionSouth
## -495.713 8.469 22.131
## regionWest log_pop18 log_poverty
## 20.054 18.008 31.676
## log_pcincome log_pop_density
## 37.688 7.655

```

```

fit_2 = lm(CRM_1000 ~ region + log_pop18 + log_poverty + log_pcincome + log_pop_density, data = pcincome)
## summary(fit_2) # aRs = 0.5355; improved with 8 significant coefs
## plot(fit_2) # rows 6, 215, 371 seem to contain outliers; treat them as influential observations...

# check collinearity and found low correlation
## check_collinearity(fit_2)

fit_3 = lm(CRM_1000 ~ (region + log_pop18 + log_poverty + log_pcincome + log_pop_density)*(region + log_pop18 + log_poverty + log_pcincome + log_pop_density))
## summary(fit_3) # aRs = 0.5748; improved, but each predictor doesn't seem as significant
fit_4 = lm(CRM_1000 ~ (region + log_pop18 + log_poverty + log_pcincome + log_pop_density)*region, data = pcincome)
## summary(fit_4) # aRs = 0.5437, 6 significant coefs
fit_5 = lm(CRM_1000 ~ (region + log_pop18 + log_poverty + log_pcincome + log_pop_density)*log_pop18, data = pcincome)
## summary(fit_5) # aRs = 0.5358, 2 significant coefs
fit_6 = lm(CRM_1000 ~ (region + log_pop18 + log_poverty + log_pcincome + log_pop_density)*log_poverty, data = pcincome)
## summary(fit_6) # aRs = 0.562, 6 significant coefs
fit_7 = lm(CRM_1000 ~ (region + log_pop18 + log_poverty + log_pcincome + log_pop_density)*log_pcincome, data = pcincome)
## summary(fit_7) # aRs = 0.5553, 4 significant coefs
fit_8 = lm(CRM_1000 ~ (region + log_pop18 + log_poverty + log_pcincome + log_pop_density)*log_pop_density, data = pcincome)
## summary(fit_8) # aRs = 0.5717, 8 significant coefs

```

We can also perform backward elimination based on BIC.

```

totalincome_model_df =
  cdi %>%
  dplyr::select(-c(pop18, pop65, bagrad, poverty, unemp, pcincome, totalinc, log_pcincome, pop_density,
lin_transform <- lm(CRM_1000 ~., data = totalincome_model_df)
bic_model5 <- stepAIC(lin_transform, direction = "backward", trace = FALSE,
                      k = log(nobs(lin_transform)))
summary(bic_model5)

##
## Call:
## lm(formula = CRM_1000 ~ region + log_bagrad + log_poverty + log_totalinc +
##     log_pop_density, data = totalincome_model_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -43.370  -9.973  -0.441   8.826 186.410
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -130.622     13.032  -10.023  < 2e-16 ***
## regionNorth Central      8.804      2.603   3.383 0.000783 ***
## regionSouth      21.386      2.576   8.304 1.31e-15 ***
## regionWest      15.723      3.222   4.880 1.50e-06 ***
## log_bagrad       9.061      3.244   2.794 0.005444 **
## log_poverty     26.620      2.078  12.810  < 2e-16 ***
## log_totalinc      6.907      1.604   4.307 2.05e-05 ***
## log_pop_density    6.038      1.252   4.823 1.96e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 18.56 on 432 degrees of freedom

```



```
## Multiple R-squared:  0.546, Adjusted R-squared:  0.5387
## F-statistic: 74.23 on 7 and 432 DF,  p-value: < 2.2e-16
```

Now, let's define our candidate models.

```
# Define model formulas and put in a list
model_A <- "CRM_1000 ~ region + log_pop_density + log_totalinc + log_bagrad + log_poverty"
model_B <- "CRM_1000 ~ region + log_pop_density + log_totalinc + log_poverty"
model_C <- "CRM_1000 ~ region + log_pop_density + log_pcincome + log_poverty"

model_list <-
  list(
    model_A = model_A,
    model_B = model_B,
    model_C = model_C
  )
```

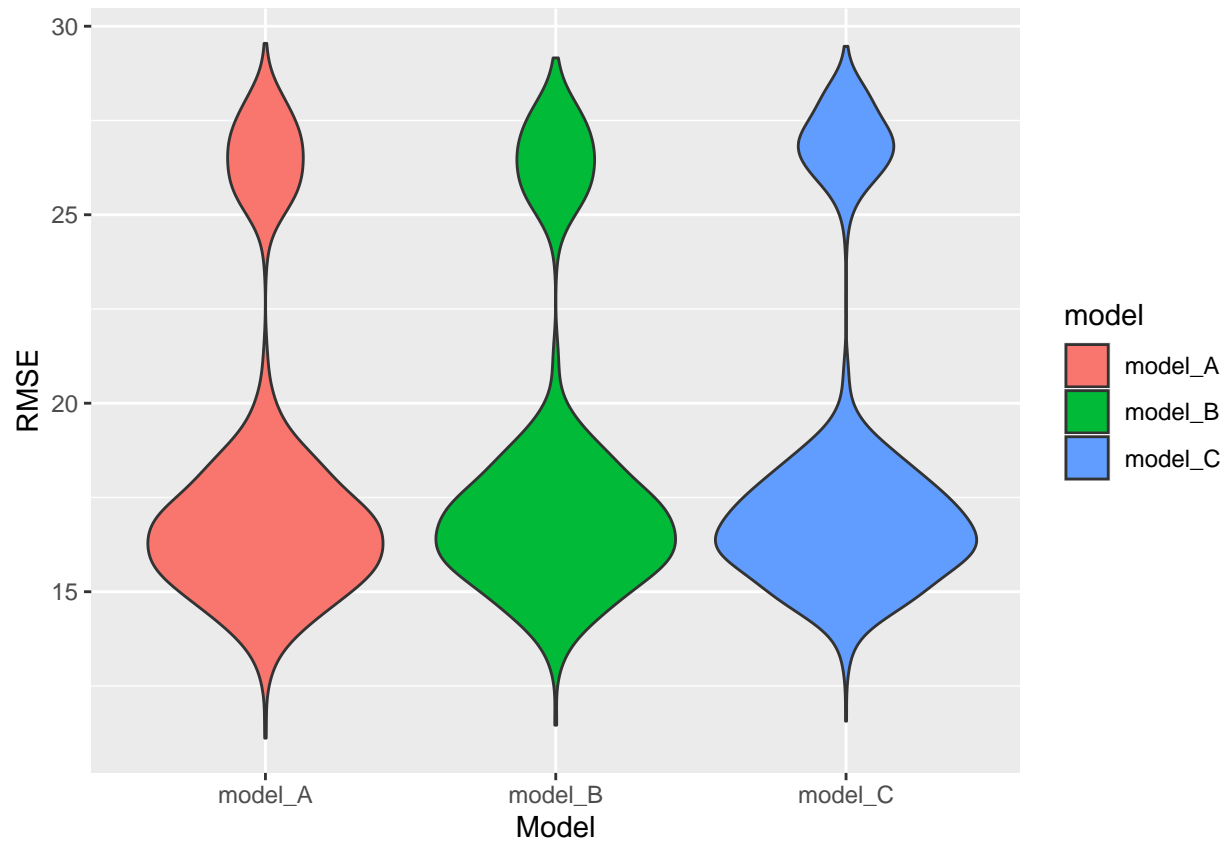
Here is each model's adjusted R-squared value.

```
# Get each model's adjusted r-squared
map(model_list, get_mod_adj_r_squared) %>%
  as_tibble() %>%
  pivot_longer(model_A:model_C,
    names_to = "model",
    values_to = "adj_r_squared") %>%
  arrange(desc(adj_r_squared)) %>%
  knitr::kable()
```

| model   | adj_r_squared |
|---------|---------------|
| model_A | 0.5386912     |
| model_B | 0.5314421     |
| model_C | 0.5283651     |

Here is the distribution of each model's cross-validation root mean squared error.

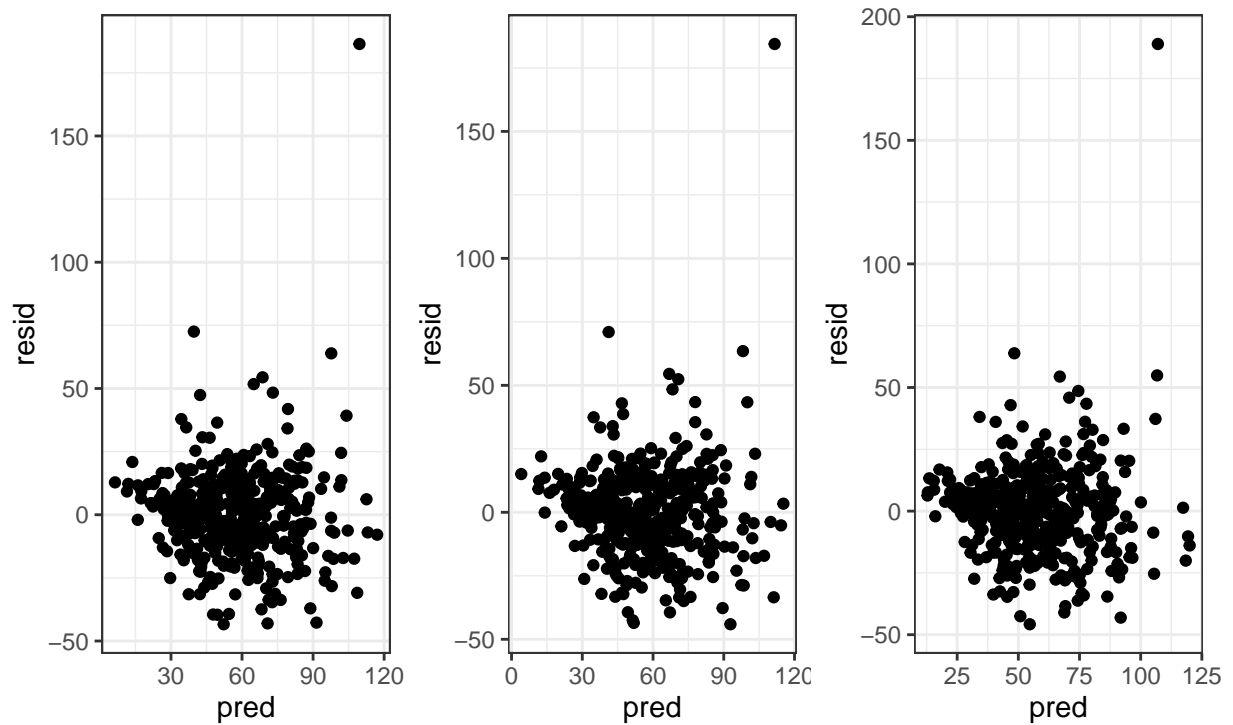
```
# Perform cross validation for each model
map(model_list, get_cv_rmse) %>%
  as_tibble() %>%
  pivot_longer(model_A:model_C,
    names_to = "model",
    values_to = "RMSE") %>%
  arrange(RMSE) %>%
  ggplot(aes(x = model, y = RMSE, fill = model)) +
  geom_violin() + labs(x = "Model", y = "RMSE")
```



Further, we can plot the model residuals as a function of the model predictions.

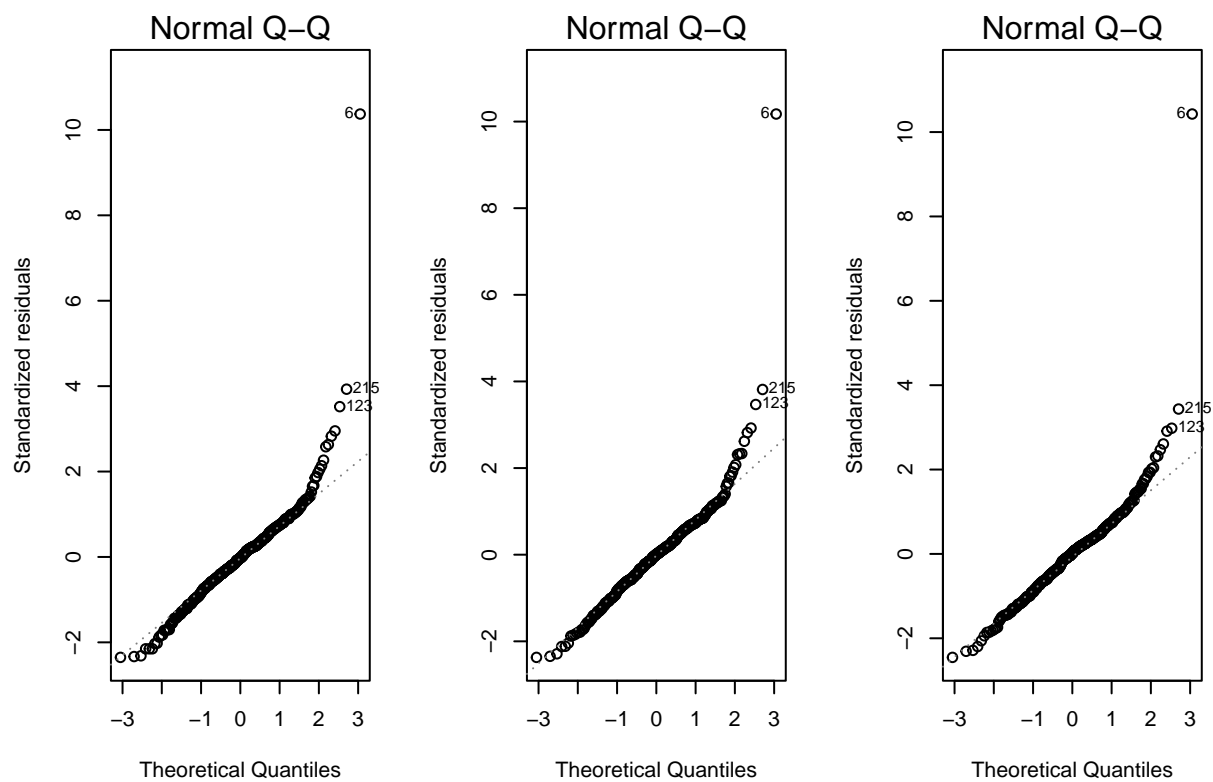
```
plot_model_residuals(model_A) +  
plot_model_residuals(model_B) +  
plot_model_residuals(model_C) +  
plot_annotation(title = 'Model Residual as a function of Model Prediction',  
                theme = theme(plot.title = element_text(size = 14)))
```

## Model Residual as a function of Model Prediction



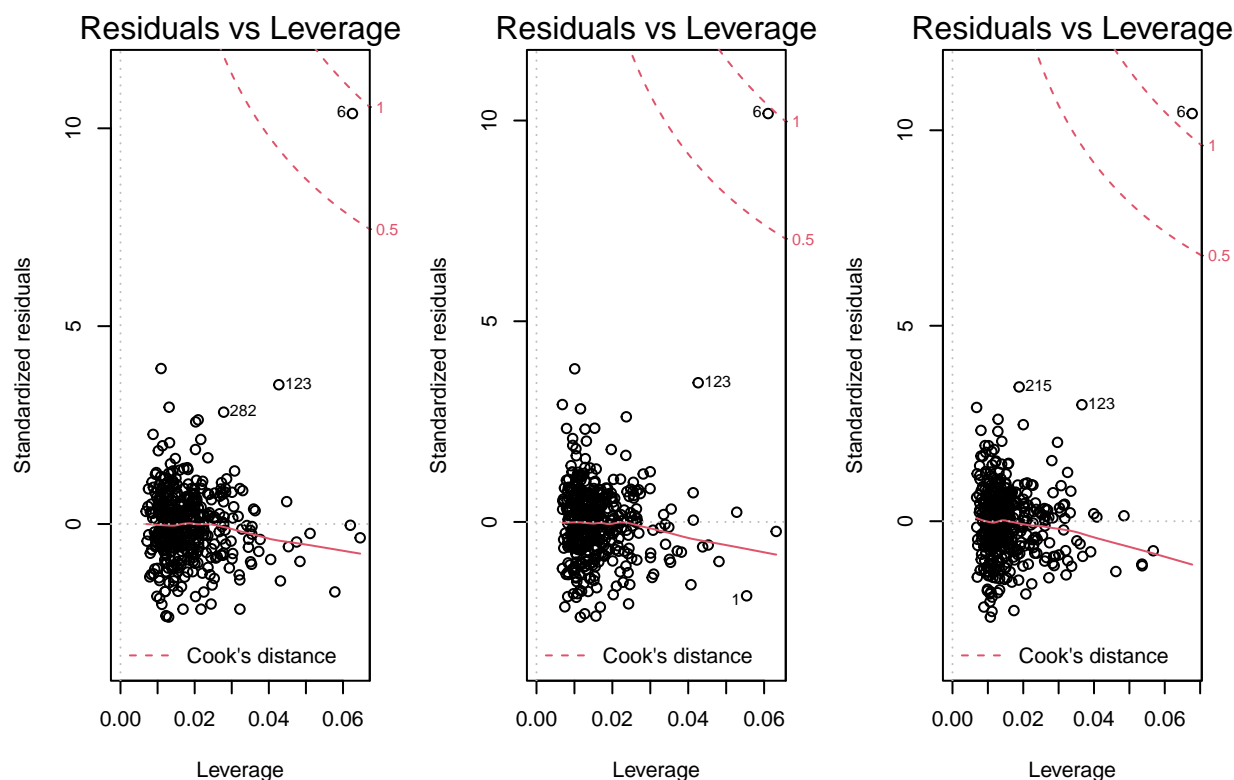
Q-Q plots

```
par(mfrow = c(1,3))
plot_mod_qq(model_A)
plot_mod_qq(model_B)
plot_mod_qq(model_C)
```



Leverage plots

```
par(mfrow = c(1,3))
plot_mod_leverage(model_A)
plot_mod_leverage(model_B)
plot_mod_leverage(model_C)
```



No wonder we had those weird distributions of RMSE in the Monte Carlo simulations - we had a serious outlier. Let's remove it and refit and re-validate the models.

```
cdi_2 <- cdi %>% slice(-6)
# Get each model's adjusted r-squared
map(model_list, get_mod_adj_r_squared, data = cdi_2) %>%
  as_tibble() %>%
  pivot_longer(model_A:model_C,
               names_to = "model",
               values_to = "adj_r_squared") %>%
  arrange(desc(adj_r_squared)) %>%
  knitr::kable()
```

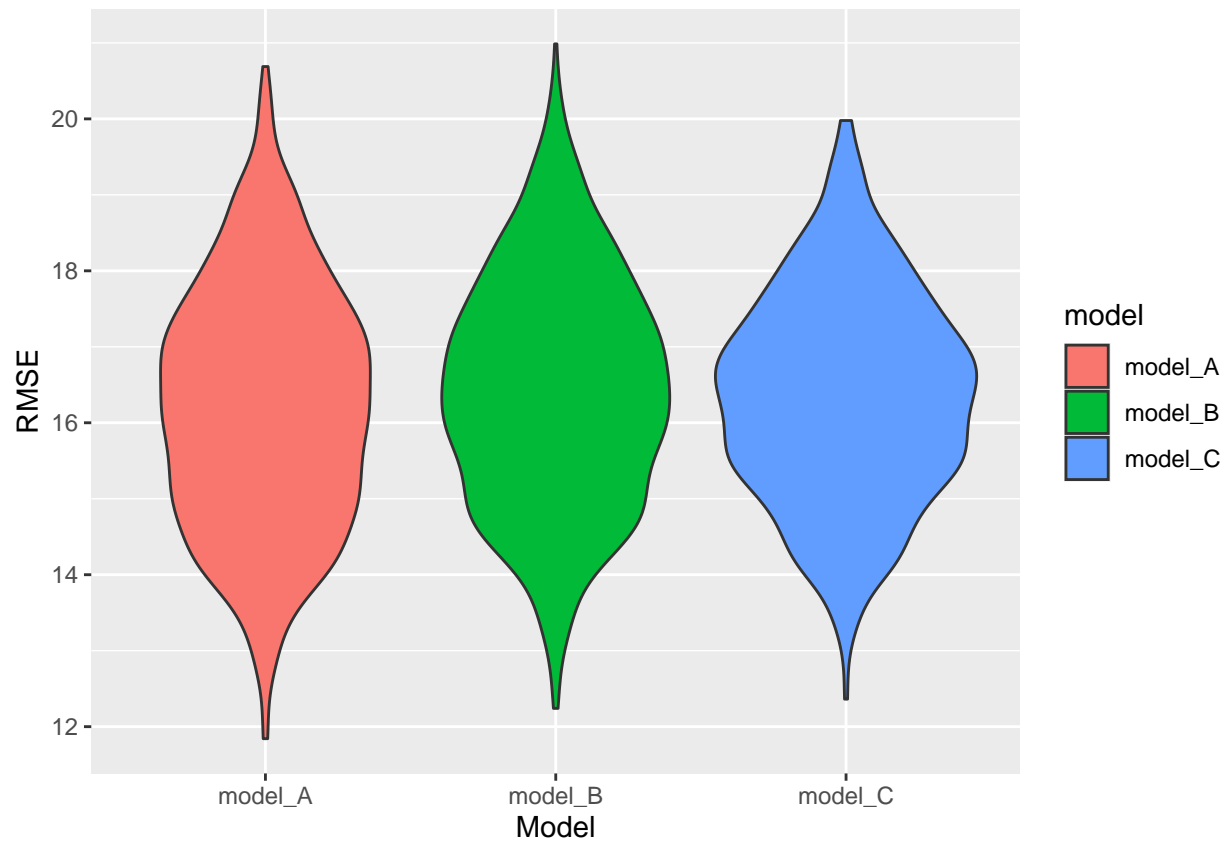
| model   | adj_r_squared |
|---------|---------------|
| model_A | 0.5804990     |
| model_C | 0.5723062     |
| model_B | 0.5682312     |

```
# Perform cross validation for each model
map(model_list, get_cv_rmse, data = cdi_2) %>%
  as_tibble() %>%
  pivot_longer(model_A:model_C,
               names_to = "model",
               values_to = "RMSE") %>%
```

```

arrange(RMSE) %>%
ggplot(aes(x = model, y = RMSE, fill = model)) +
geom_violin() + labs(x = "Model", y = "RMSE")

```

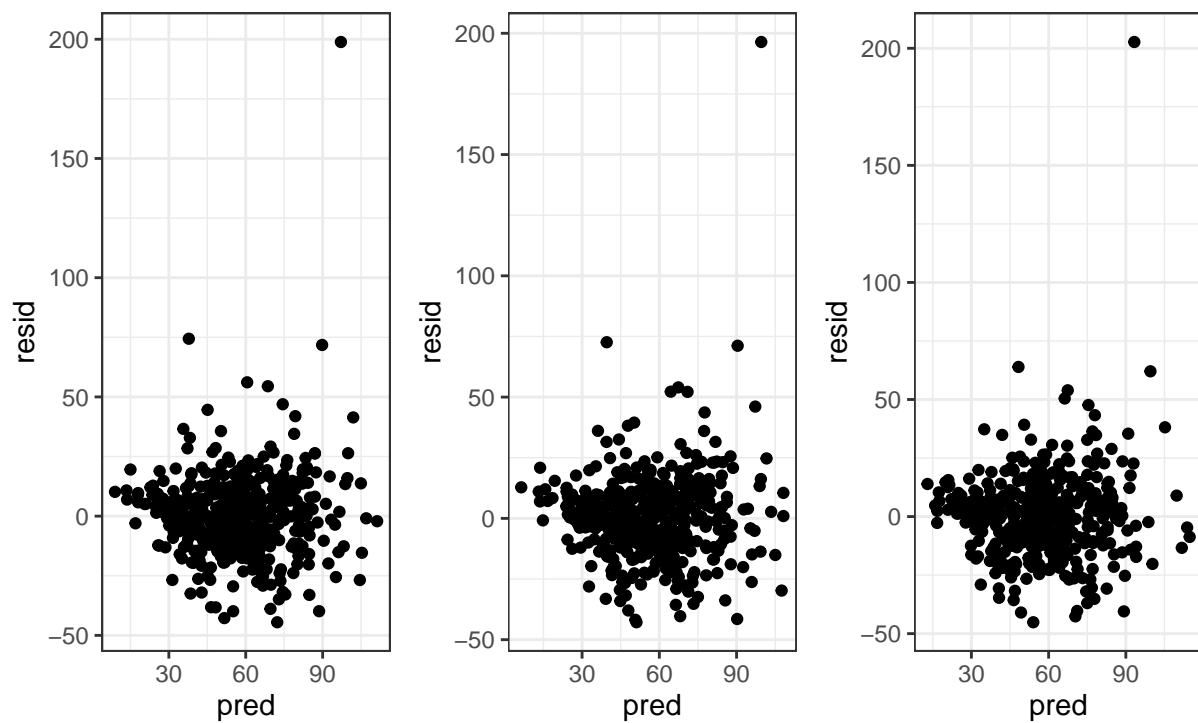


```

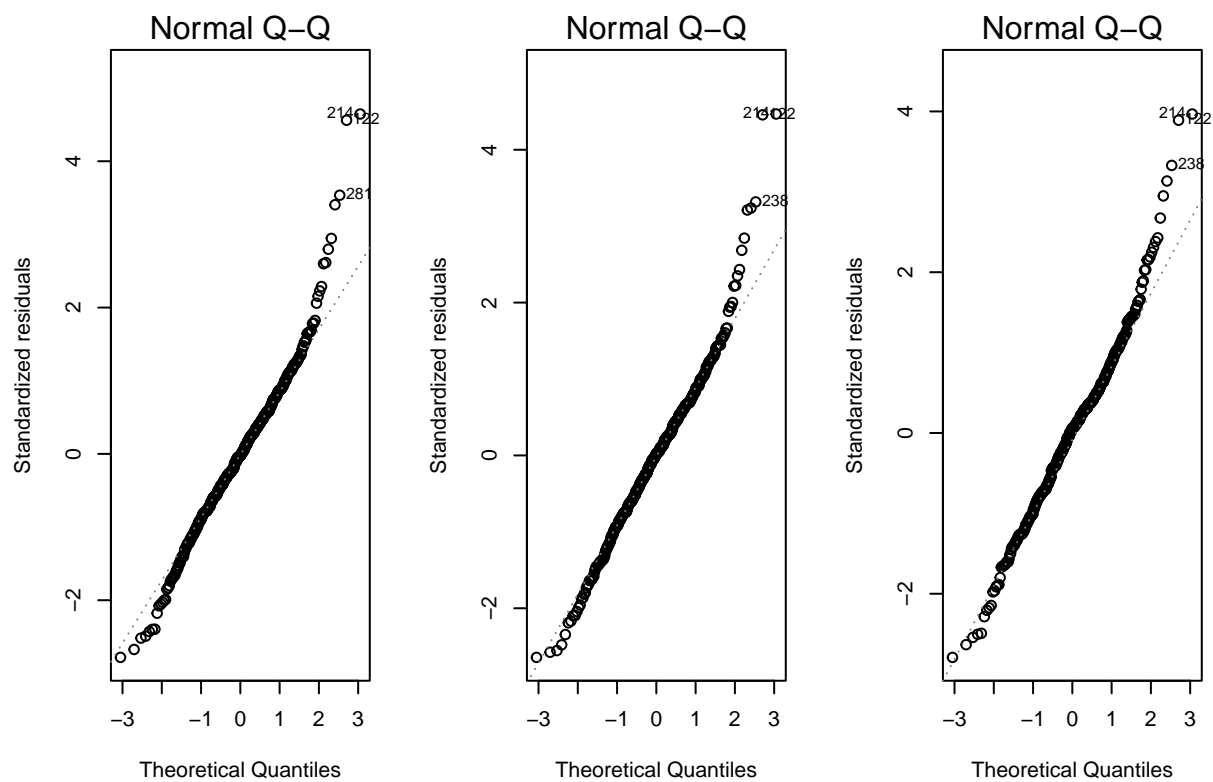
#Residual plots
plot_model_residuals(model_A, data = cdi_2) +
plot_model_residuals(model_B, data = cdi_2) +
plot_model_residuals(model_C, data = cdi_2) +
plot_annotation(title = 'Model Residual as a function of Model Prediction',
  theme = theme(plot.title = element_text(size = 14)))

```

## Model Residual as a function of Model Prediction



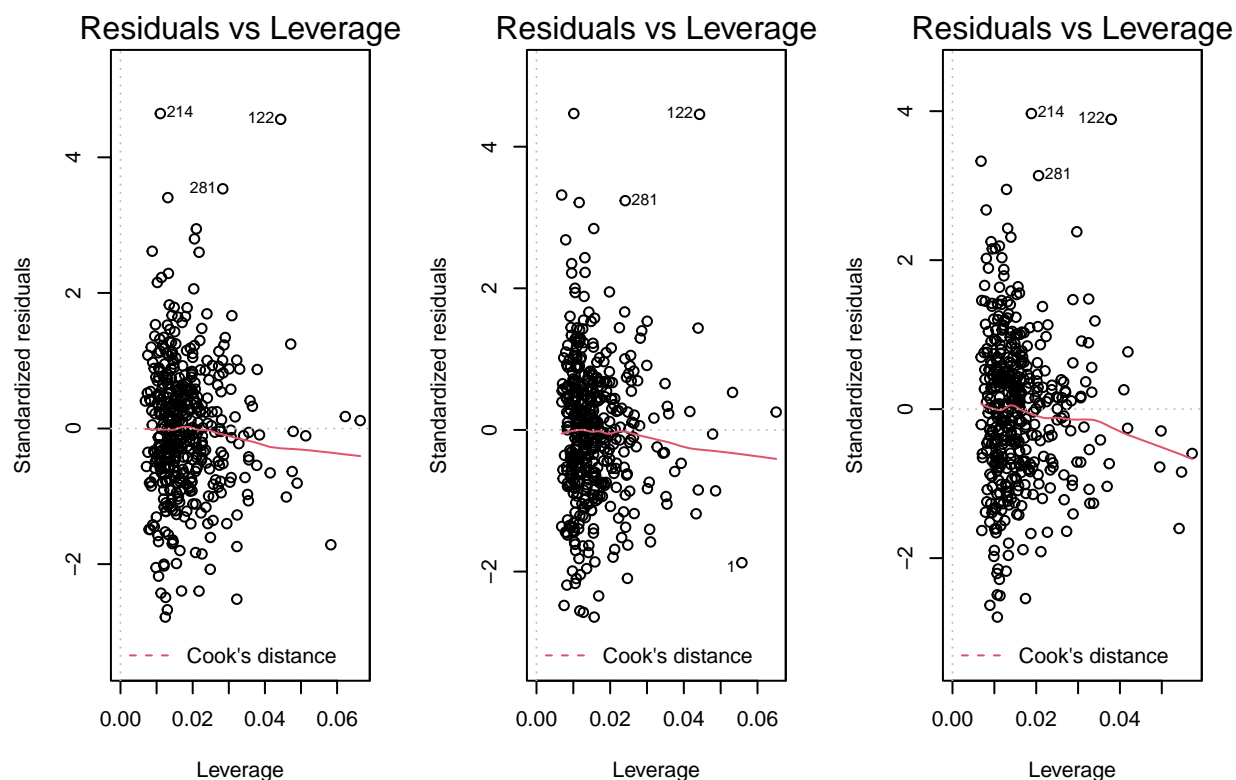
```
#Q-Q plots  
par(mfrow = c(1,3))  
plot_mod_qq(model_A, data = cdi_2)  
plot_mod_qq(model_B, data = cdi_2)  
plot_mod_qq(model_C, data = cdi_2)
```



#### *#Leverage plots*

```
plot_mod_leverage(model_A, data = cdi_2)
plot_mod_leverage(model_B, data = cdi_2)
plot_mod_leverage(model_C, data = cdi_2)
```





For reasons listed in the paper, we will use model B. Let's summarize the model.

```
broom::tidy(lm(model_B, data = cdi_2)) %>%
  knitr::kable()
```

| term                | estimate   | std.error | statistic  | p.value  |
|---------------------|------------|-----------|------------|----------|
| (Intercept)         | -98.510093 | 9.181019  | -10.729756 | 0.00e+00 |
| regionNorth Central | 11.058492  | 2.298291  | 4.811615   | 2.10e-06 |
| regionSouth         | 25.572291  | 2.235019  | 11.441643  | 0.00e+00 |
| regionWest          | 18.224954  | 2.786970  | 6.539343   | 0.00e+00 |
| log_pop_density     | 4.555372   | 1.104612  | 4.123956   | 4.47e-05 |
| log_totalinc        | 8.316132   | 1.386287  | 5.998855   | 0.00e+00 |
| log_poverty         | 21.188450  | 1.605139  | 13.200384  | 0.00e+00 |